



Shared memory on the GPU

Jesper Mosegaard

Plan

- Shared Memory on the GPU
 - PBuffer
 - FBO
- Memory addressing
- Memory write
 - MRT



Shared memory on the GPU


■ Problem

- ☐ How do fragment communicate ?
- ☐ How can calculation depend on previous calculations ?

■ Solution

- ☐ “Render-to-texture”
- ☐ Reading from texture

Shared Memory



Solutions to rtt

- `glCopyTexImage2D`
 - Copy of current frame-buffer to a texture
 - Slow...
- PBuffer and `renderTarget` based
 - Render to offscreen Pixel buffer, bind buffer as texture.
 - Faster, but each PBuffer has an OpenGL context
 - WGL extension
- Framebuffer Objects
 - Attaching texture images to framebuffer objects
 - Not dependant on WGL extensions
 - Potentially faster than PBuffers



PBuffer Material

<http://www.ati.com/developer/ATIpbuffer.pdf>

http://developer.nvidia.com/object/gdc_oglrtt.html

http://oss.sgi.com/projects/ogl-sample/registry/ARB/wgl_pbuffer.txt

http://oss.sgi.com/projects/ogl-sample/registry/ARB/wgl_render_texture.txt

<http://www.gpgpu.org/developer/>

http://www.opengl.org/resources/features/GL_EXT_render_target.txt

http://www.opengl.org/discussion_boards/cgi_directory/ultimatebb.cgi?ubb=get_topic;f=3;t=011406

RTT through PBuffer

- `WGL_ARB_pbuffer`
 - Pixel buffers are additional non-visible rendering buffers for an OpenGL renderer
- `WGL_ARB_render_texture`
 - Allows a color buffer to be used for both rendering and texturing
- A way to implement render-to-texture
 - (WGL = on windows only)

Creating the PBuffer

- Query device and rendering contexts
 - *wglGetCurrentDC()* *wglGetCurrentContext()*
- Choose pixel format supporting bind to texture
 - *wglChoosePixelFormatARB()*
- Create the PBuffer
 - *wglCreatePbufferARB()*

Rendering to the Pbuffer

- Switch rendering context, frame buffer writes.
 - *wglMakeContextCurentARB()*

Binding PBuffer as texture

- Create Texture Object
 - *glGenTextures()*
- Bind Texture Object
 - *glBindTexture()*
- Bind PBuffer to currently bound texture object
 - *wglBindTexImageARB()*
- Release PBuffer
 - *wglReleaseTexImageARB()*

PBuffer problems

■ Pbuffers Give us Render-To-Texture

- Designed to create an environment map or two
- Never intended to be used for GPGPU (100s of pbuffers)

□ Problem

- Each pbuffer has its own OpenGL render context
- Each pbuffer may have depth and/or stencil buffer
- Changing OpenGL contexts is slow

□ Solution

- Many optimizations to avoid this bottleneck...
- Allocate several surfaces

PBuffer Tips

■ Ping-Pong

- Create a PBuffer with a front and back buffer
 - Switch between the two, rendering and reading
 - Avoids context switch
- Use auxillary buffers for more surfaces



Floating point textures, material

http://oss.sgi.com/projects/ogl-sample/registry/ATI/texture_float.txt

http://www.nvidia.com/dev_content/nvopenglspecs/GL_NV_float_buffer.txt

Floating point textures

- Vendor Specific Extensions
 - ATI_texture_float (Also supported by NV40)
 - NV_float_buffer
- General
 - ARB_texture_float (only on NV)
- Support for filtering may be limited / implemented in software


Key Differences

- ARB_texture_float
 - 16bit and 32bit textures

- Older:
 - ATI_texture_float
 - Available in Fixed Function and Programs
 - NV_float_buffer
 - Rectangle Textures
 - Available ONLY in Programs

Floating point texture formats

- FP32 (IEEE standard)
 - s23e8 (largest counting number: 16,777,217)
- FP16
 - s10e5 (largest : 2,049)
- Remember, the ATI Radeon 9800 can ONLY work with 24 bit internally
 - s16e7 (largest : 131,073)



$$\text{sign} * 1.\text{mantissa} * 2^{(\text{exponent} + \text{bias})}$$

Using PBuffers for your projects

- PBuffers have been wrapped up nicely in the RenderTexture library – see RenderTexture.h
- Construct new RenderTexture
 - `rt = new RenderTexture();`
- Setup format
 - `rt->Reset("rgba=32f tex2D rtt");`
- Bind for rendering
 - `rt->BeginCapture();`
- Bind to the current active texture
 - `glActiveTextureARB(GL_TEXTURE0_ARB);`
 - `rt->BindBuffer(WGL_FRONT_LEFT_ARB);`

Framebuffer Object

- EXT_framebuffer_object
- Advantages
 - Single OpenGL context
 - Switching framebuffers is potentially faster than switching OpenGL context (*wglMakeCurrent*)
 - Share renderbuffer images and texture images between framebuffers
 - Easier to use.

Framebuffer Object

- OpenGL Framebuffer, logical buffers
 - Color, depth, stencil, accumulation
- EXT_framebuffer_object
 - Attach and detach texture and renderBuffers as logical framebuffer objects



Framebuffer Object, material

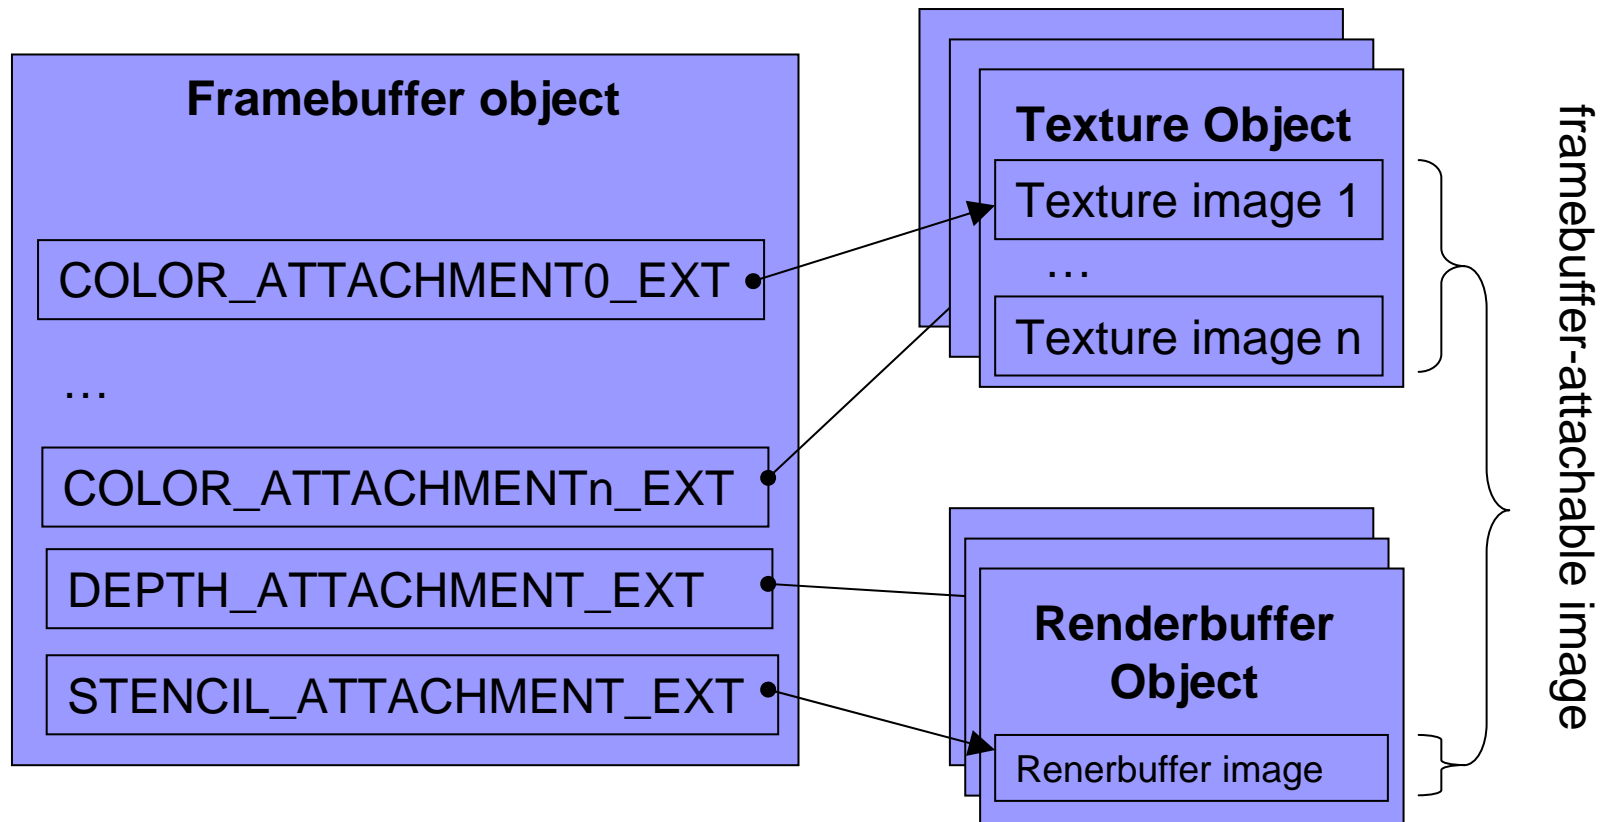
http://www.opengl.org/documentation/extensions/EXT_framebuffer_object.txt

http://download.nvidia.com/developer/presentations/2005/GDC/OpenGL_Day/OpenGL_FrameBuffer_Object.pdf

Framebuffer_object

- Application-generated framebuffer object
 - Collection framebuffer-attachable image
 - Offscreen buffers (renderbuffer)
 - Textures (texture image)

Object architecture



New Terminology

- *Framebuffer*
 - Collection of framebuffer attachable images
 - State (direction of GL rendering)
 - Windows-system framebuffer and application-generated
- *Renderbuffer* – offscreen image and state
 - *Renderbuffer image* – 2D array of pixels. Part of a renderbuffer.
- *Framebuffer-attachable image* – 2D array of pixels that can be attached to a framebuffer. (e.g. texture image and renderbuffer image)
- *Attachement point* – state that references a framebuffer-attachable image in a framebuffer. (color, depth, stencil)
- *Attach* – connecting one object to another. Similar to bind.

FBO: Create framebuffer

```
glGenFramebuffersEXT(); // generate a framebuffer
```

```
glGenTextures(); // generate a texture
```

```
// bind the generated framebuffer
```

```
glBindFramebufferEXT(GL_FRAMEBUFFER_EXT, framebufferID);
```

```
glBindTexture(GL_TEXTURE_2D, textureID); // setup texture parameters
```

```
glTexImage2D(GL_TEXTURE_2D, 0, GL_RGBA32F_ARB, w, h, 0, GL_RGBA,  
             GL_FLOAT, NULL);
```

```
// Attaching the two Texture Images to the Framebuffer
```

```
glFramebufferTexture2DEXT(GL_FRAMEBUFFER_EXT,  
                          GL_COLOR_ATTACHMENT0_EXT, GL_TEXTURE_2D, textureID, 0);
```


FBO: Render to image of framebuffer

// Bind application-generated framebuffer ID

glBindFramebufferEXT(GL_FRAMEBUFFER_EXT, ID);

glDrawBuffer(GL_COLOR_ATTACHMENT0_EXT); // render attachment 0

DRAW OPENGGL STUFF

// Bind window-system framebuffer object

glBindFramebufferEXT(GL_FRAMEBUFFER_EXT, 0);

FBO: Bind Texture

- Just like usual !

```
glEnable(GL_TEXTURE_2D);  
glBindTexture(GL_TEXTURE_2D, TEXID);
```



Performance tips


- Don't create and destroy FBO's every frame
- Try to avoid modifying textures used as rendering destinations using TexImage, CopyTexImage etc.

FBO Usage Scenarios

- FBO allows several ways of switching between rendering destinations. In order of increasing performance:
 - Multiple FBOs
 - create a separate FBO for each texture you want to render to
 - switch using `BindFramebuffer()`
 - can be 2x faster than `wglMakeCurrent()` in beta NVIDIA drivers
 - Single FBO, multiple texture attachments
 - textures should have same format and dimensions
 - use `FramebufferTexture()` to switch between textures
 - Single FBO, multiple texture attachments
 - attach textures to different color attachments
 - use `glDrawBuffer ()` to switch rendering to different color attachments

Memory addressing

- CPU defined vertex attributes
 - Define few vertices
- Vertex based calculations
- Interpolated by rasterizer
 - Exploit this if possible
- Fragment based calculations
 - Creates dependent texture lookups (problem on radeon 9800)





Texture lookup

- Tip:

- Move as much independent math after a texture fetch to hide the latency

Warning

■ Floating-point can leave unaddressable texels

- Nvidia FP32 (16,777,217)
- Nvidia FP16 (2,049)
- Ati 24-bit float: (131,073)



The tradeoff of parameters

- Draw quad with four texture coordinates
 - Fast drawing, interpolation not flexible
- Draw pixel sized points with individual texture coordinates
 - Very slow draw, very flexible

Memory write

- Output from fragment program
 - Only one 4-vector (rgba)
- Multi Render Target (MRT)
 - Write to several "surfaces" (or framebuffer attached images)



MRT: Material

http://oss.sgi.com/projects/ogl-sample/registry/ATI/draw_buffers.txt

MRT

- Render to several (up to four at the moment) surfaces at the same position
- OpenGL:
 - void **DrawBuffersATI**(sizei n, const enum *bufs);
- Fragment program
 - Specify: ATI_draw_buffers option
 - Use: result.color[n]
- Cg
 - COLOR1 through COLOR3 binding semantics